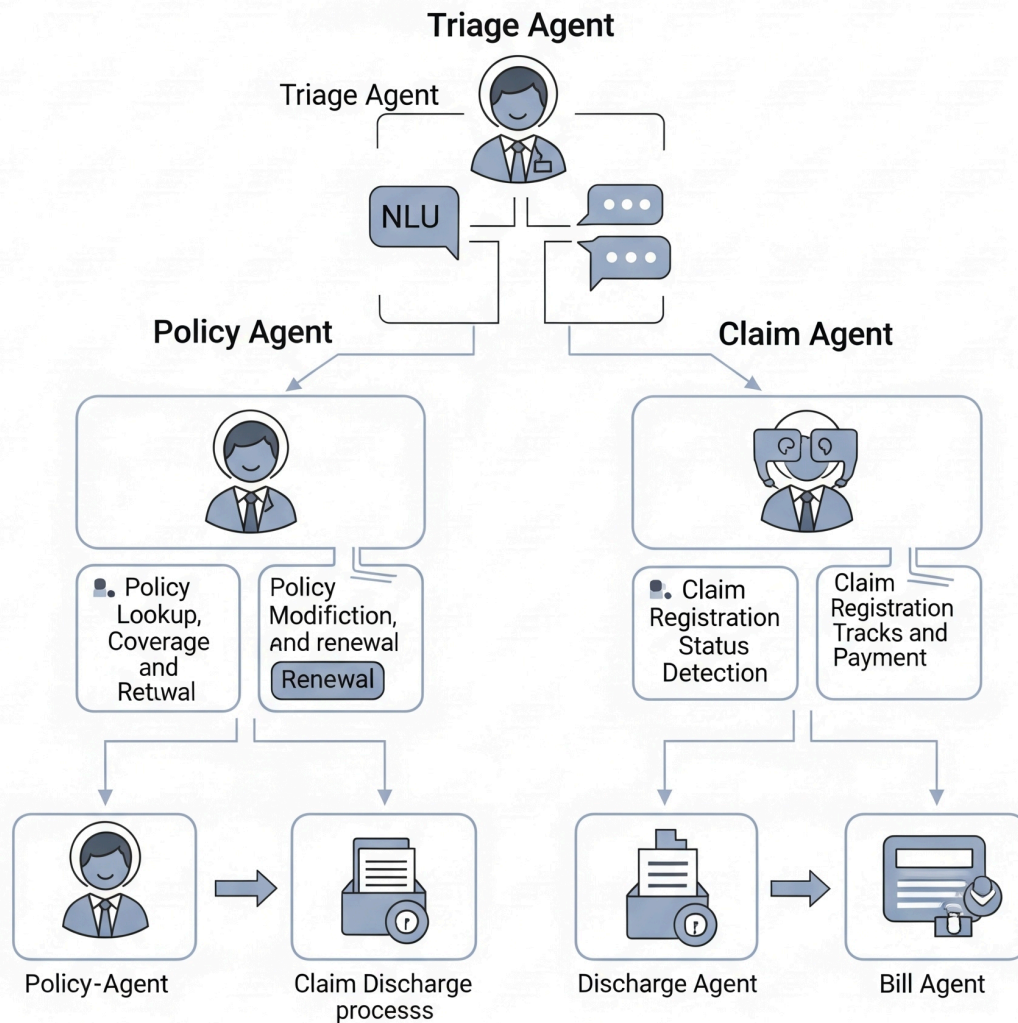


AI Agent System for Customer Service



1. Overall Architecture Overview

The system is designed as a cloud-native, serverless, and highly scalable application. It follows a microservices-oriented approach where the core ADK agent logic is encapsulated within a containerized service deployed on Google Cloud Run. This service interacts with a managed PostgreSQL database for persistent data storage.

2. Component Breakdown

2.1. Google Cloud Run

Google Cloud Run serves as the primary hosting environment for the ADK agent.

- **Serverless Container Platform:** Cloud Run allows deploying containerized applications without managing the underlying infrastructure. It automatically scales up or down based on incoming request traffic, even scaling to zero instances when idle, which optimizes costs.
- **Request-Driven:** It's ideal for stateless, request-driven services like a web API, ensuring that the ADK agent is always ready to receive and process user requests.
- **Built-in Features:** Provides automatic HTTPS, custom domain mapping, traffic splitting, and integration with other Google Cloud services.

2.2. Google ADK + FastAPI Web Server

FastAPI is chosen as the web framework within the Cloud Run container.

- **High Performance:** Built on Starlette and Pydantic, FastAPI offers excellent performance comparable to Node.js and Go.
- **Asynchronous Support:** Its native support for `async/await` is crucial for handling concurrent requests efficiently, especially when the ADK agent might be waiting on external API calls or database operations.

2.3. Google ADK Agentic Framework

The Google ADK (Agent Development Kit) framework powers the core intelligence and orchestration of the agent.

- **Agentic Capabilities:** Provides the foundational components for building AI agents, including capabilities for understanding natural language, managing conversation state, reasoning, and tool utilization.
- **Tool Integration:** Facilitates the integration of various tools (e.g., `available_policies`, `purchase_policy`, `check_ongoing_claim`, etc).
- These tools use Fire Store for storing state along with session state provided by ADK.
- **Orchestration:** Manages the flow of interaction, deciding which tools to use and when, based on user input and prompt instructions.

3. Data Flow and Interactions

1. **User Request:** An end-user initiates an interaction (e.g., via a chat interface or an application) that sends an HTTP request to the ADK agent.
2. **Cloud Run Invocation:** The request triggers an instance of the ADK Agent service on Google Cloud Run.
3. **ADK's + FastAPI Server:** FastAPI receives the HTTP request and directs it to the appropriate endpoint.
4. **ADK Framework Execution:** The FastAPI endpoint invokes the Google ADK agentic framework. The framework processes the user's input, leveraging an LLM (e.g., Gemini 2.0 Flash) for natural language understanding and determining the appropriate action.
5. **Tool Execution:** Based on the agent's reasoning, it may invoke one or more internal "tools."
 - **Database Interaction:** If a tool requires data (e.g., `check_existing_policy`, `check_ongoing_claim`), the ADK framework, via FastAPI, connects to Google Firestore (Firebase) to query or update the necessary information.
 - **ADK Session Store:** Google ADK uses SQL data store for persistence powered by PostgreSQL.
6. **Response Generation:** After executing tools and processing data, the ADK framework generates a natural language response (if applicable) or a structured data response.
7. **User Receives Response:** The response is returned to the end-user and rendered by Google's ADK web UI.

Google ADK Multi Agent Workflow Architecture

